



# ARP protocol: same LAN

❖ A wants to send datagram to B

- B's MAC address not in A's ARP table.

❖ A **broadcasts** ARP query packet, containing B's IP address

- dest MAC address = FF-FF-FF-FF-FF-FF
- all nodes on LAN receive ARP query

❖ B receives ARP packet, replies to A with its (B's) MAC address

- frame sent to A's MAC address (unicast)

- A caches (saves) IP-to-MAC address pair in its ARP table until information becomes old (times out)

- soft state: information that times out (goes away) unless refreshed


- ARP is “plug-and-play”:

- nodes create their ARP tables *without intervention from net administrator*

# Addressing: routing to another LAN


walkthrough: send datagram from A to B via R

- focus on addressing – at IP (datagram) and MAC layer (frame)
- assume A knows B's IP address
- assume A knows IP address of first hop router, R (how?)
- assume A knows R's MAC address (how?)




# Addressing: routing to another LAN

- ❖ A creates IP datagram with IP source A, destination B
- ❖ A creates link-layer frame with R's MAC address as dest, frame contains A-to-B IP datagram




# Addressing: routing to another LAN

- ❖ frame sent from A to R
- ❖ frame received at R, datagram removed, passed up to IP




# Addressing: routing to another LAN

- ❖ R forwards datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram



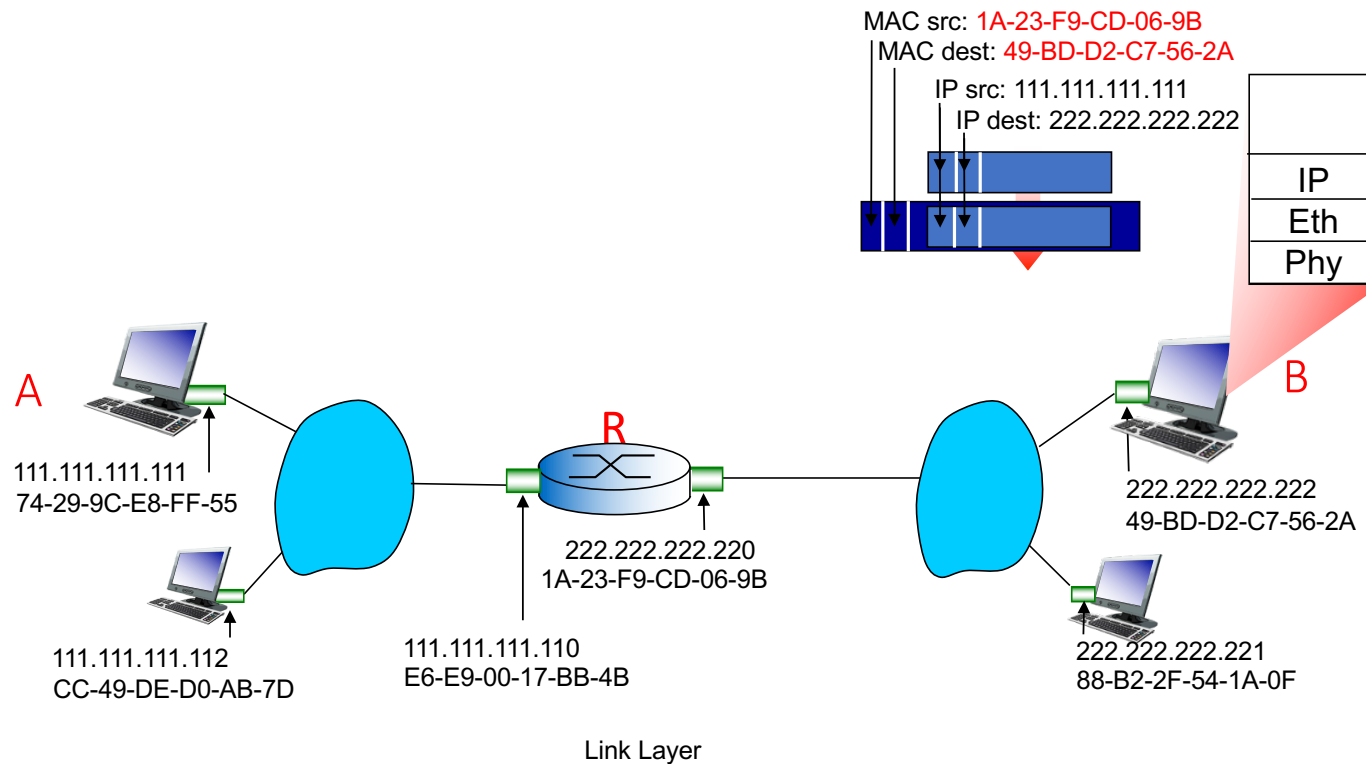
# Addressing: routing to another LAN

- ❖ R forwards datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram



# Addressing: routing to another LAN

- ❖ R forwards datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram





# Link layer, LANs: outline

5.1 introduction, services

5.2 error detection,  
correction

5.3 multiple access  
protocols

## 5.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

5.5 link virtualization:  
MPLS


5.6 data center  
networking

5.7 a day in the life of a  
web request

# Ethernet

“dominant” wired LAN technology:

- cheap <\$20 for NIC
- first widely used LAN technology
- simpler, cheaper than token LANs and ATM
- kept up with speed race: 10 Mbps – 10 Gbps




Metcalfe's Ethernet sketch


Link Layer

# Ethernet: physical topology

- **bus**: popular through mid 90s
  - all nodes in same collision domain (can collide with each other)
- **star**: prevails today
  - active **switch** in center
  - each “spoke” runs a (separate) Ethernet protocol (nodes do not collide with each other)




**bus**: coaxial cable



Link Layer

# Ethernet frame structure

sending adapter encapsulates IP datagram (or other network layer protocol packet) in **Ethernet frame**



## *preamble:*

- ❖ 7 bytes with pattern 10101010 followed by one byte with pattern 10101011
- ❖ used to synchronize receiver, sender clock rates

## Ethernet frame structure (more)

❖ *addresses*: 6 byte source, destination MAC addresses

- if adapter receives frame with matching destination address, or with broadcast address (e.g. ARP packet), it passes data in frame to network layer protocol
- otherwise, adapter discards frame

❖ *type*: indicates higher layer protocol (mostly IP but others possible, e.g., Novell IPX, AppleTalk)

❖ *CRC*: cyclic redundancy check at receiver

- error detected: frame is dropped



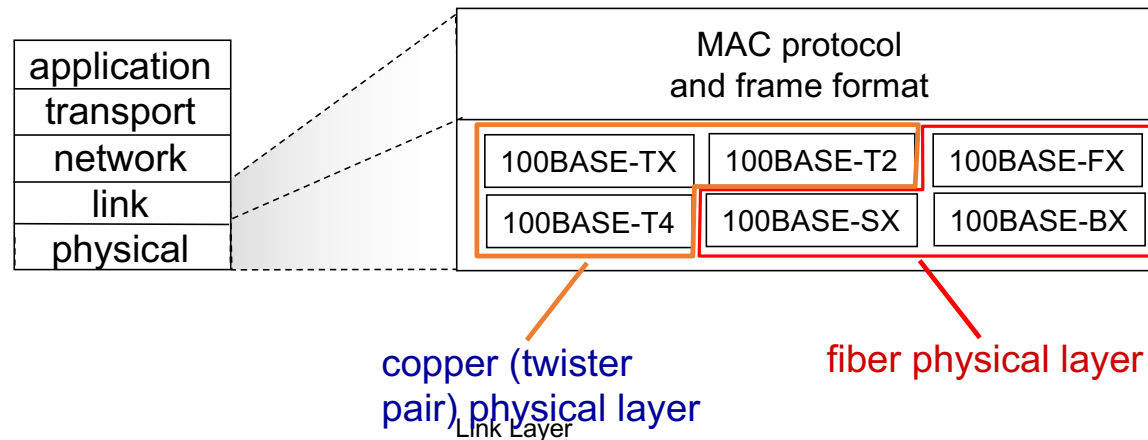
# Ethernet: unreliable, connectionless

- ❖ *connectionless*: no handshaking between sending and receiving NICs
- ❖ *unreliable*: receiving NIC doesn't send acks or nacks to sending NIC
  - data in dropped frames recovered only if initial sender uses higher layer rdt (e.g., TCP), otherwise dropped data lost
- ❖ Ethernet's MAC protocol: unslotted *CSMA/CD with binary backoff*

## 802.3 Ethernet standards: link & physical layers

### ❖ *many* different Ethernet standards

- common MAC protocol and frame format
- different speeds: 2 Mbps, 10 Mbps, 100 Mbps, 1Gbps, 10G bps
- different physical layer media: fiber, cable



# Link layer, LANs: outline

5.1 introduction, services

5.2 error detection,  
correction

5.3 multiple access  
protocols

## 5.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

5.5 link virtualization:  
MPLS

5.6 data center  
networking

5.7 a day in the life of a  
web request



# Ethernet switch

## ❖ link-layer device: takes an *active* role

- store, forward Ethernet frames
- examine incoming frame's MAC address, *selectively* forward frame to one-or-more outgoing links when frame is to be forwarded on segment, uses CSMA/CD to access segment

## ❖ *transparent*


- hosts are unaware of presence of switches

## ❖ *plug-and-play, self-learning*

- switches do not need to be configured

## Switch: *multiple* simultaneous transmissions

- ❖ hosts have dedicated, direct connection to switch
- ❖ switches buffer packets
- ❖ Ethernet protocol used on *each* incoming link, but no collisions; full duplex
  - each link is its own collision domain
- ❖ *switching*: A-to-A' and B-to-B' can transmit simultaneously, without collisions



## Switch forwarding table


Q: how does switch know A' reachable via interface 4, B' reachable via interface 5?

❖ A: each switch has a **switch table**, each entry:

- (MAC address of host, interface to reach host, time stamp)
- looks like a routing table!

Q: how are entries created, maintained in switch table?

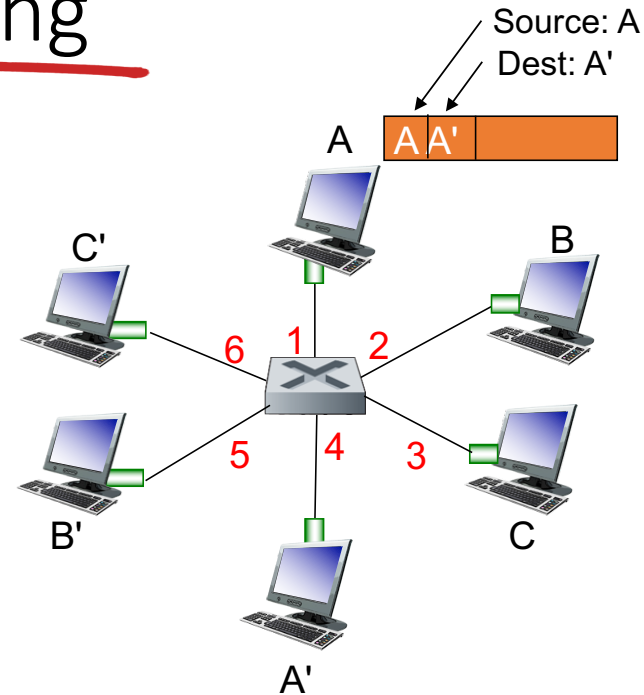
- something like a routing protocol?



switch with six interfaces  
(1,2,3,4,5,6)

# Switch: self-learning

- switch *learns* which hosts can be reached through which interfaces
  - when frame received, switch “learns” location of sender: incoming LAN segment
  - records sender/location pair in switch table



MAC addr	interface	TTL
A	1	60

Switch table  
(initially empty)

Link Layer

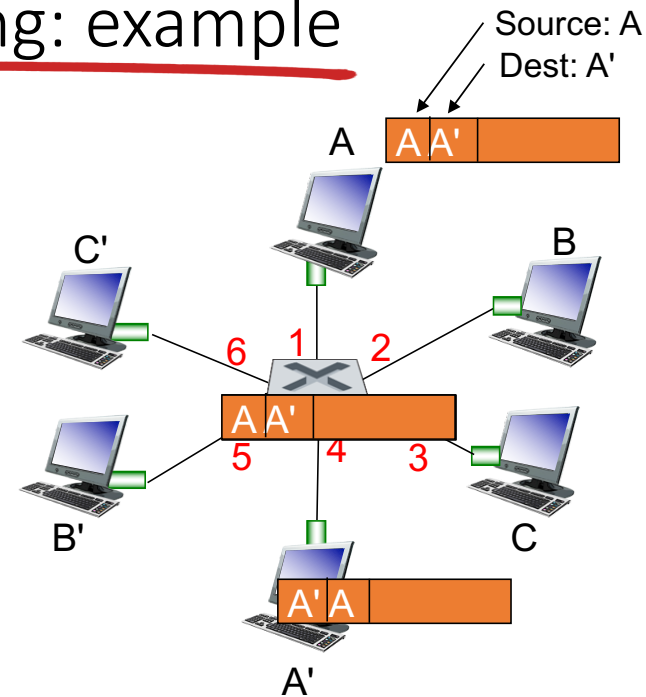
# Switch: frame filtering/forwarding

when frame received at switch:

1. record incoming link, MAC address of sending host
2. index switch table using MAC destination address
3. **if** entry found for destination  
    **then** {  
        **if** destination on segment from which frame arrived  
            **then** drop frame  
            **else** forward frame on interface indicated by entry  
    }  
    **else** flood /\* forward on all interfaces except arriving  
                interface \*/

## Self-learning, forwarding: example

- ❖ frame destination, A', location unknown: **flood**
- ❖ destination A location known: **selectively send on just one link**




MAC addr	interface	TTL
A	1	60
A'	4	60

switch table  
(initially empty)

Link Layer

# Interconnecting switches


❖ switches can be connected together



**Q:** sending from A to G - how does  $S_1$  know to forward frame destined to F via  $S_4$  and  $S_3$ ?

❖ **A:** self learning! (works exactly the same as in single-switch case!)

# ~~Institutional network~~



Link Layer



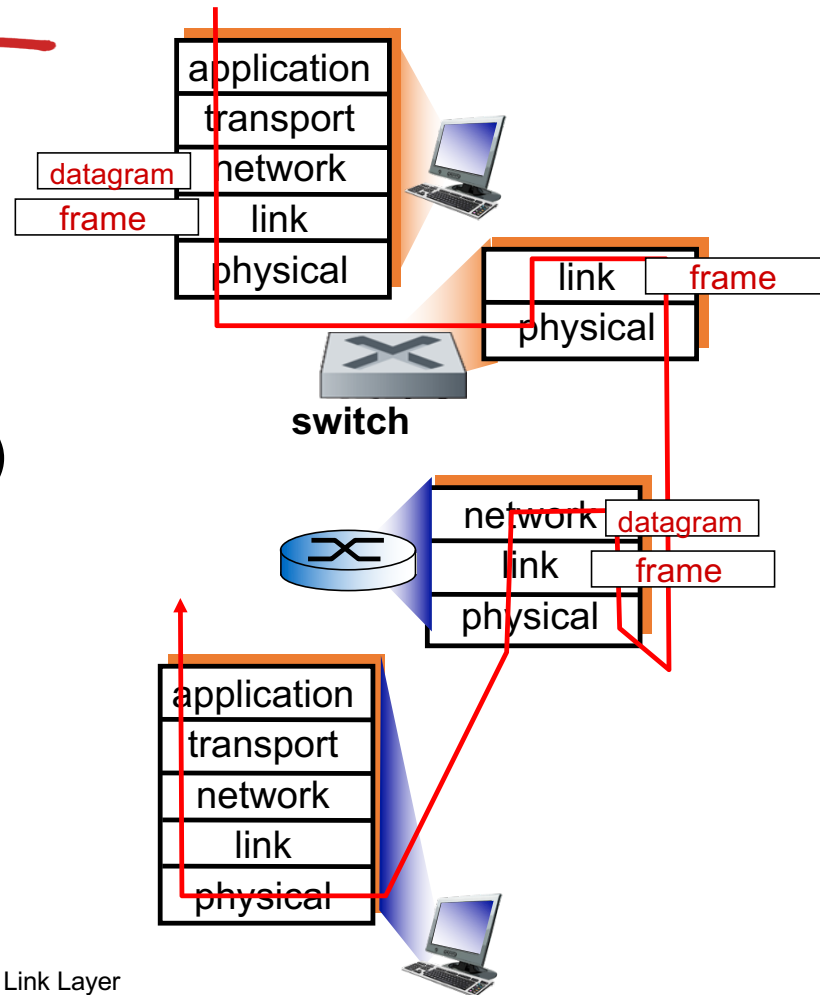
# Switches vs. routers

both are store-and-forward:

- **routers:** network-layer devices (examine network-layer headers)
- **switches:** link-layer devices (examine link-layer headers)


both have forwarding tables:

- **routers:** compute tables using routing algorithms, IP addresses
- **switches:** learn forwarding table using flooding, learning, MAC addresses



Link Layer

# VLANs: motivation



## *consider:*


- ❖ CS user moves office to EE, but wants connect to CS switch?
- ❖ single broadcast domain:
  - all layer-2 broadcast traffic (ARP, DHCP, unknown location of destination MAC address) must cross entire LAN
  - security/privacy, efficiency issues

# VLANs

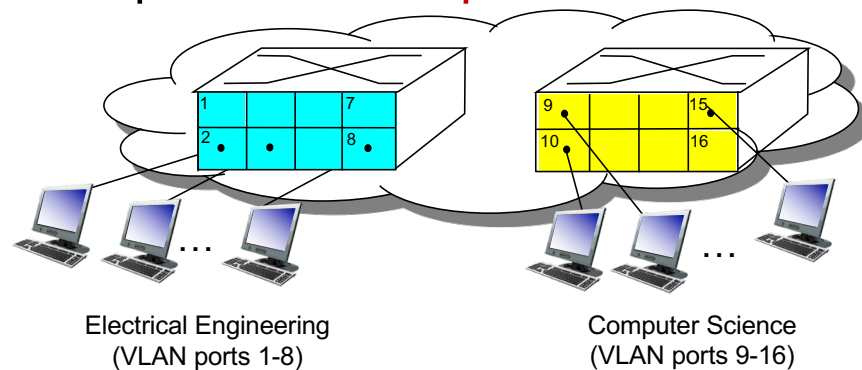
## Virtual Local Area Network

switch(es) supporting VLAN capabilities can be configured to define multiple **virtual** LANS over single physical LAN infrastructure.

**port-based VLAN:** switch ports grouped (by switch management software) so that **single** physical switch .....



... operates as **multiple** virtual switches



Link Layer

# Port-based VLAN


❖ **traffic isolation:** frames to/from ports 1-8 can *only* reach ports 1-8

- can also define VLAN based on MAC addresses of endpoints, rather than switch port


❖ **dynamic membership:** ports can be dynamically assigned among VLANs

❖ **forwarding between VLANs:** done via routing (just as with separate switches)

- in practice vendors sell combined switches plus routers



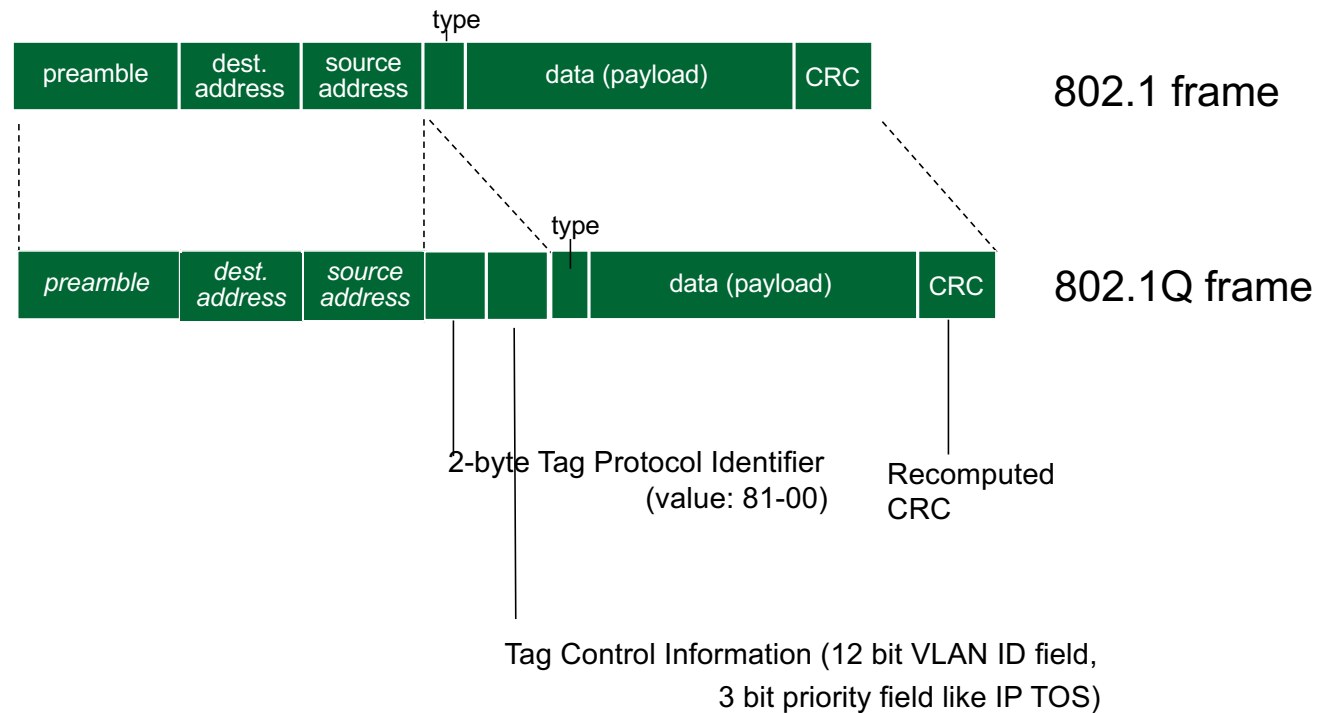
# ~~VLANs spanning multiple switches~~



❖ *trunk port*: carries frames between VLANs defined over multiple physical switches

- frames forwarded within VLAN between switches can't be vanilla 802.1 frames (must carry VLAN ID info)
- 802.1q protocol adds/removed additional header fields for frames forwarded between trunk ports

# 802.1Q VLAN frame format



# Link layer, LANs: outline

- 5.1 introduction, services
- 5.2 error detection, correction
- 5.3 multiple access protocols
- 5.4 LANs
  - addressing, ARP
  - Ethernet
  - switches
  - VLANS
- 5.5 link virtualization: MPLS
- 5.6 data center networking
- 5.7 a day in the life of a web request

# Data center networks

- ❖ 10's to 100's of thousands of hosts, often closely coupled, in close proximity:
  - e-business (e.g. Amazon)
  - content-servers (e.g., YouTube, Akamai, Apple, Microsoft)
  - search engines, data mining (e.g., Google)

- ❖ challenges:
  - multiple applications, each serving massive numbers of clients
  - managing/balancing load, avoiding processing, networking, data bottlenecks




Inside a 40-ft Microsoft container,  
Chicago data center



# Data center networks


## load balancer: application-layer routing

- receives external client requests
- directs workload within data center
- returns results to external client (hiding data center internals from client)



# Data center networks

- ❖ rich interconnection among switches, racks:
  - increased throughput between racks (multiple routing paths possible)
  - increased reliability via redundancy



# Link layer, LANs: outline

- 5.1 introduction, services
- 5.2 error detection, correction
- 5.3 multiple access protocols
- 5.4 LANs
  - addressing, ARP
  - Ethernet
  - switches
  - VLANS
- 5.5 link virtualization: MPLS
- 5.6 data center networking
- 5.7 a day in the life of a web request

## Synthesis: a day in the life of a web request

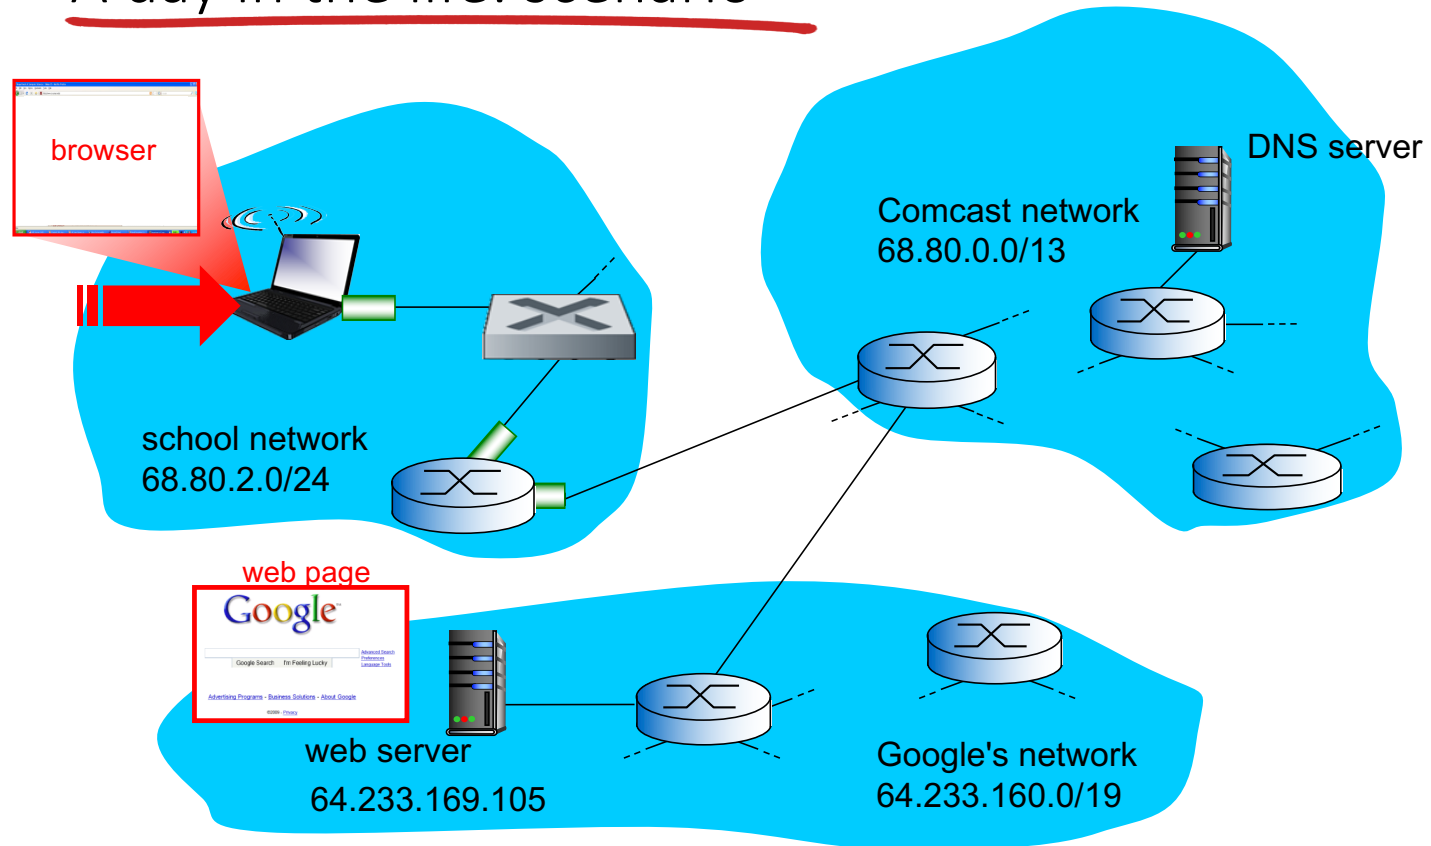
### ❖ journey down protocol stack complete!

- application, transport, network, link


### ❖ putting-it-all-together: synthesis!

- *goal*: identify, review, understand protocols (at all layers) involved in seemingly simple scenario: requesting www page
- *scenario*: student attaches laptop to campus network, requests/receives www.google.com

## A day in the life: scenario




## A day in the life... connecting to the Internet



- ❖ connecting laptop needs to get its own IP address, addr of first-hop router, addr of DNS server: use **DHCP**
- ❖ DHCP request **encapsulated** in **UDP**, encapsulated in **IP**, encapsulated in **802.3** Ethernet
- ❖ Ethernet frame **broadcast** (dest: FFFFFFFFFFFFFFFF) on LAN, received at router running **DHCP** server
- ❖ Ethernet **demuxed** to IP demuxed, UDP demuxed to DHCP


## A day in the life... connecting to the Internet



- ❖ DHCP server formulates *DHCP ACK* containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- ❖ encapsulation at DHCP server, frame forwarded (*switch learning*) through LAN, demultiplexing at client
- ❖ DHCP client receives DHCP ACK reply

Client now has IP address, knows name & addr of DNS server, IP address of its first-hop router

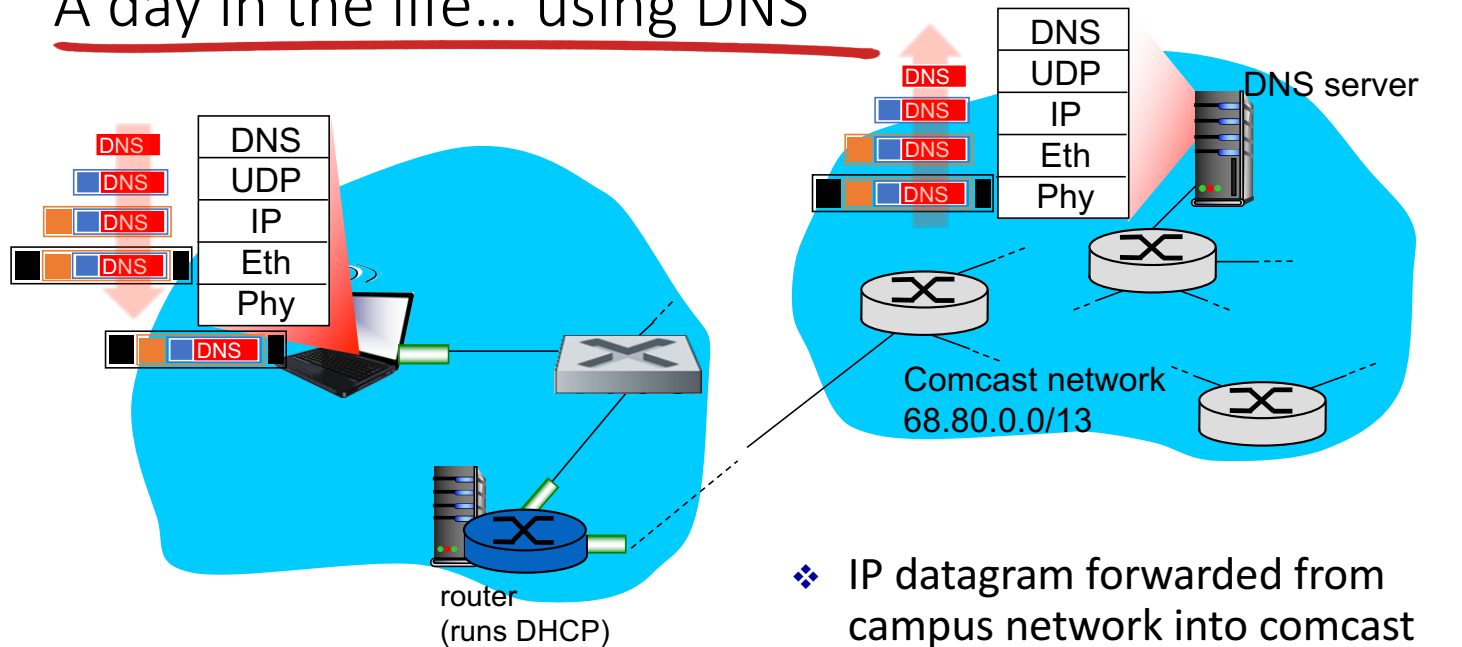
## A day in the life... ARP (before DNS, before HTTP)



- ❖ before sending **HTTP** request, need IP address of `www.google.com`: **DNS**
- ❖ DNS query created, encapsulated in UDP, encapsulated in IP, encapsulated in Eth. To send frame to router, need MAC address of router interface: **ARP**
- ❖ **ARP query** broadcast, received by router, which replies with **ARP reply** giving MAC address of router interface
- ❖ client now knows MAC address of first hop router, so can now send frame containing DNS query



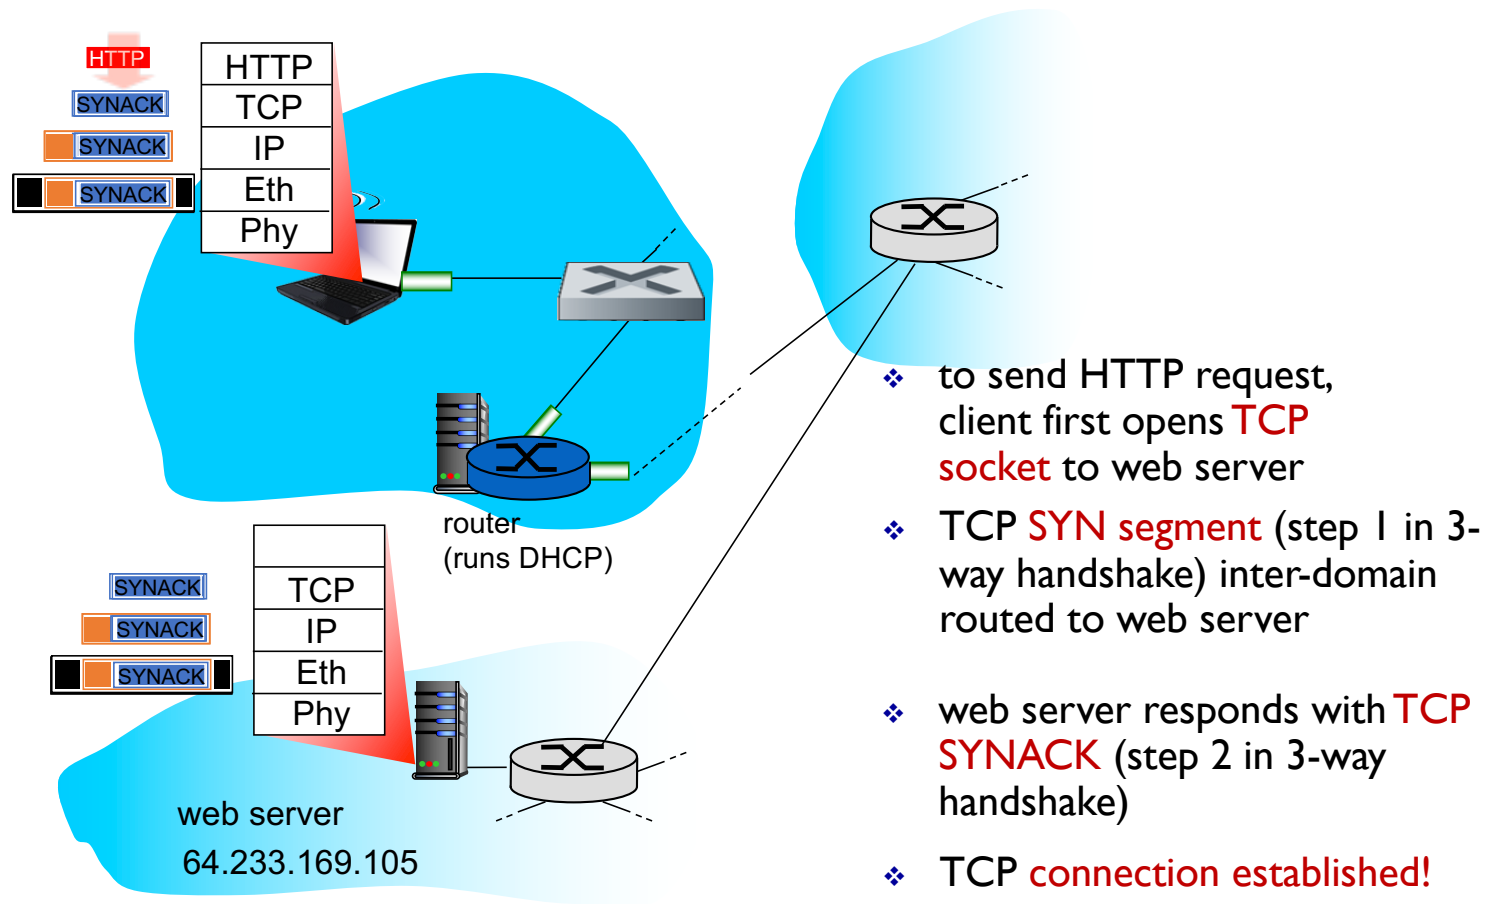
## A day in the life... using DNS



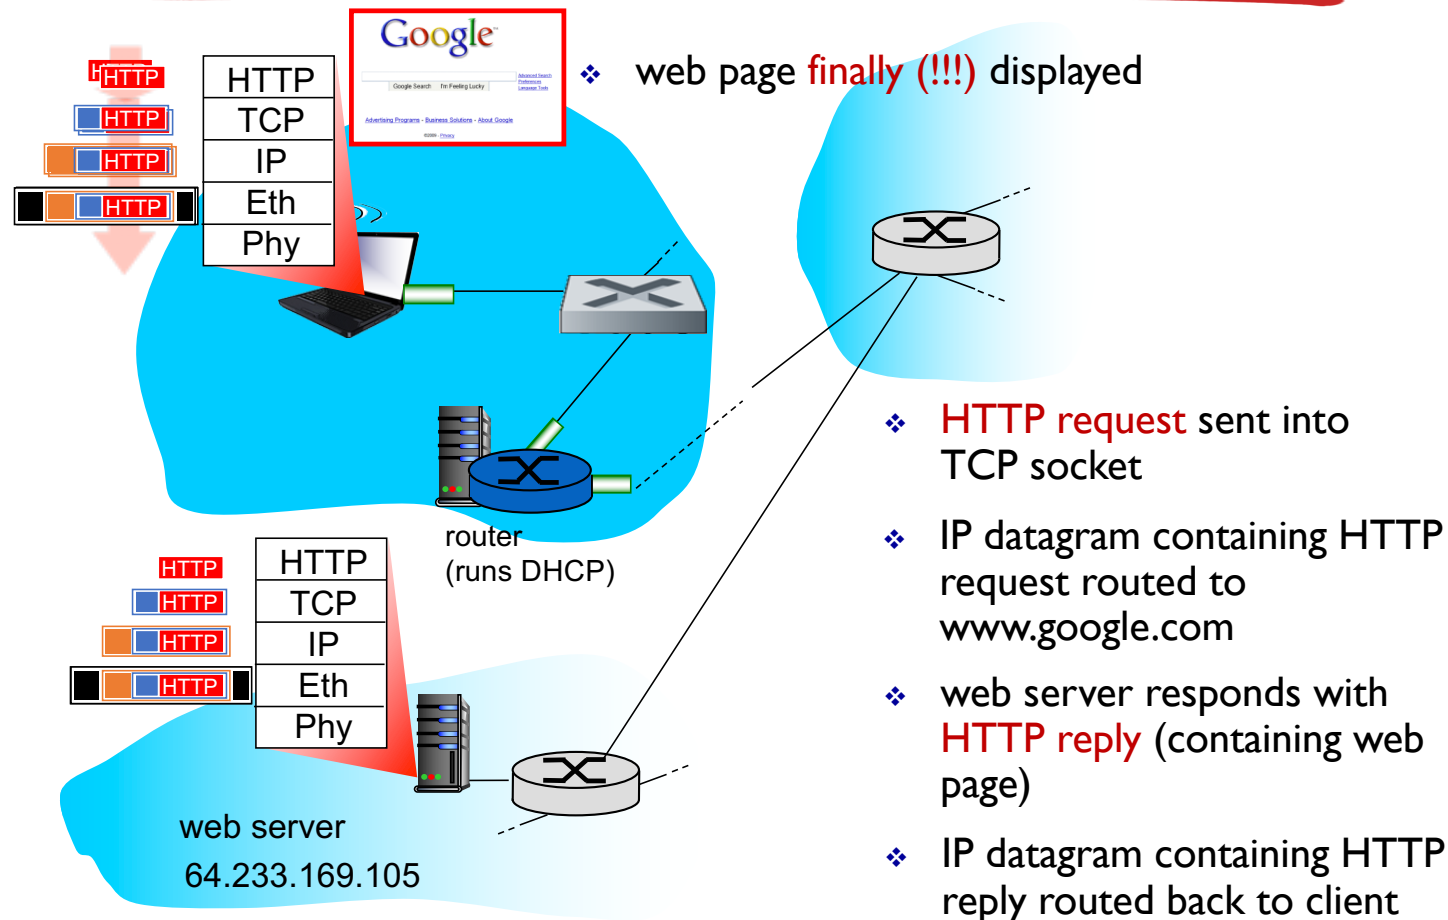
- ❖ IP datagram containing DNS query forwarded via LAN switch from client to 1<sup>st</sup> hop router

- ❖ IP datagram forwarded from campus network into comcast network, routed (tables created by **RIP**, **OSPF**, **IS-IS** and/or **BGP** routing protocols) to DNS server
- ❖ demux'ed to DNS server
- ❖ DNS server replies to client with IP address of **www.google.com**

## A day in the life...TCP connection carrying HTTP



# A day in the life... HTTP request/reply



Link Layer